

Approximating Data Movement via Compiler Support

Junhan Zhou
Vignesh Balaji


The Problem

Data movement is a major bottleneck for performance scaling of parallel programs

Cores communicate through:


Synchronization

The blue locks ensures consistent updates to data values by multiple threads



Cache coherence


The arrows shows data movement between cores via cache coherence



Potential Solution

Data movement can be reduced by writing better code (smaller critical regions, better synchronization strategies, etc.)


OR we could use **Approximate Computing!**




People approximated **synchronization** before...

But **data movement** still exists for coherence...

So we try approximating cache **coherence!**
(Incoherence cache)



Heuristic 1: Estimating **importance** of Program variables



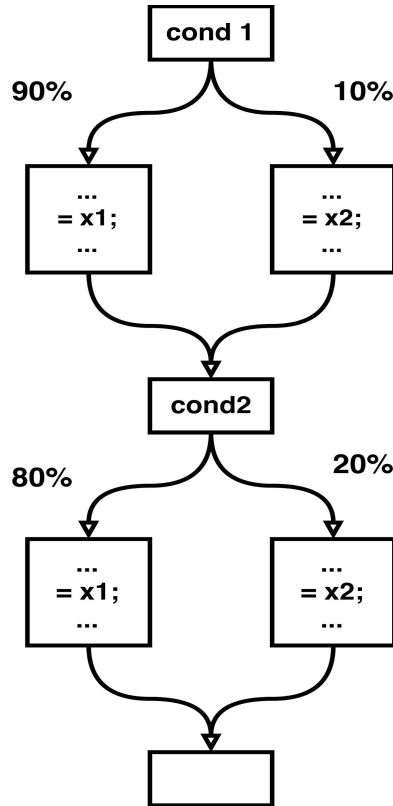
- Count the **number** of uses of variables that are *shared* among different threads
- Variables that are used less might be better approximation targets (lead to lesser quality degradation)
- Knob:

Uses Threshold = 1



Uses
Threshold =
n

Heuristic 2: **Quality impact** of approximating variables




➤ Track if a variable is present on all **paths** at every meet point

➤ Variables that appear only on few paths are not likely to cause significant quality degradation


➤ Knob:

Path freq. < 0%




Path freq. < 100%

Heuristic 3: Detecting **imbalance** in computation



- Detect the **amount** of computation on variables in different threads
- Variables that are heavily operated upon only in a few threads do not need coherence for *all* threads
- Knob:

Uses Threshold = 1



Uses Threshold = n


Implementation Details

First identify **potential shared variables** to approximate (global variables with load-modify-store pattern inside locks that does not determine control flow)

Then use the mentioned **heuristics** to **select** the shared variable to approximate.

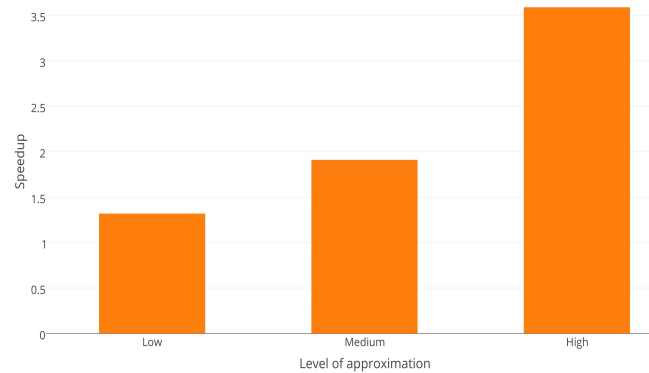
Simulated the **performance** effect of incoherent caches by not adding cycles for coherence actions (directory lookup, invalidations, etc.)

For **error** numbers, manipulated thread local copies of data followed by an approximate merge of values




Results

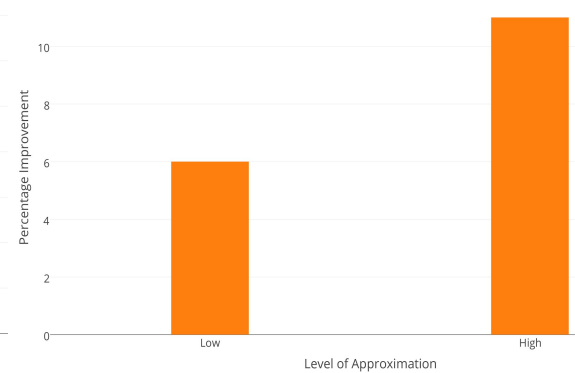
H1 Performance Plot



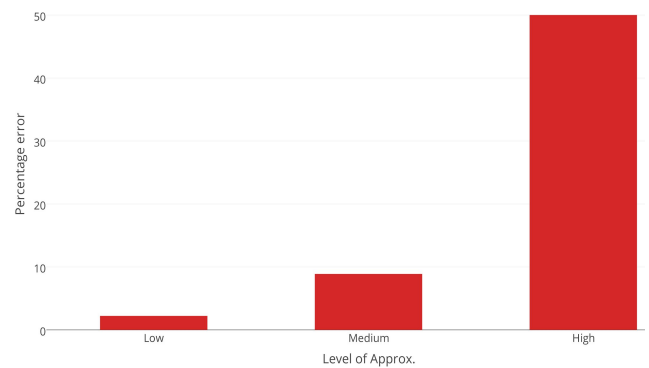
H2 Performance Plot



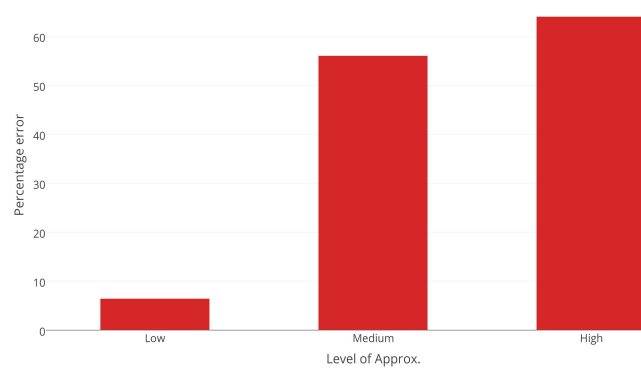
H3 Performance Plot



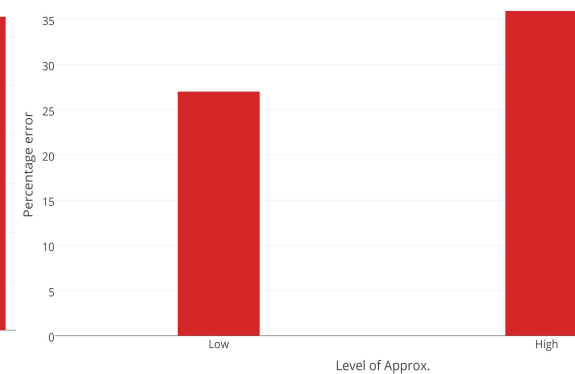
H1 error plot



H2 error plot



H3 error plot



Limitations and Future Work

Our pass can be improved by incorporating:

- Code annotations (where the programmer identifies approximable parts of code [1][2][3])
- Feedback from code profiling to get an idea on the dynamic aspects of the code [3]

[1] Adrian Sampson, Werner Dietl, Emily Fortuna, Danushen Gnanapragasam, Luis Ceze, and Dan Grossman. 2011. EnerJ: approximate data types for safe and general low-power computation. In *Proceedings of the 32nd ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI '11)*

[2] Sampson, Adrian, et al. "Accept: A programmer-guided compiler framework for practical approximate computing." *University of Washington Technical Report UW-CSE-15-01 1* (2015)

[3] Sasa Misailovic, Michael Carbin, Sara Achour, Zichao Qi, and Martin C. Rinard. 2014. Chisel: reliability- and accuracy-aware optimization of approximate computational kernels. In *Proceedings of the 2014 ACM International Conference on Object Oriented Programming Systems Languages & Applications (OOPSLA '14)*